CodeBandhu

Bridging Legal Literacy Gaps with AI-Powered Summarization and Simplification

Devyansh Nigam
B.E-CSE
Chandigarh University
Mohali, Punjab, India
nigamdevyansh90@gmail.com

Er. Geetkiran Kaur
CSE
Chandigarh University
Mohali, Punjab, India
geetkiran.cse@cumail.in

Gautam Kumar
B.E-CSE
Chandigarh University
Mohali, Punjab, India
23BCS13028@cuchd.in

Abstract—The legal documents issued by the Government of India tend to be formal, lengthy, opaque, and supremely difficult for the ordinary citizen to digest. While parsing this legalese is a daunting task for any one person, we take the position that a citizen cannot have enough cognitive ability to extract accurate meaning and therefore take action on their own behalf or for their community. This paper describes an AI-driven LegalTech framework called CodeBandhu that takes any Indian legal or policy document, and simplifies, summarizes, and suggests actions therefrom surfaces as recommend actions for users.

The CodeBandhu system draws upon NLP (Natural Language Processing) techniques for unsupervised text simplification, abstractive and extractive summarization, and argument-mining to automatically take heirarchical illegal dense material, and present in simplified and nuggeted form.

The model's modular pipeline consists of document normalisation, lexical simplification using Legal-BERT for basic simplification, rhetorical role annotations, and a recommender engine based on legal issues and back precedents recognized. CodeBandhu is trained and evaluated against publicly available datasets, including notably LawSum and IN-Abs (Indian Abstractive Legal Summaries), and validation along both automatic metrics (ROUGE, BLEU, F1-Intent) as well as generated for experts and then validated by human experts. The results demonstrate an enormously increase in readability and relevance of the summaries, and an alignment with legal practitioners preferences. The aim of CodeBandhu is to assist scraping and simplifying legal, policy the materials in a citizen friendly format and suggesting feasible actions to promote greater legal literacy, as well as faster compreh veering to and civic acting on corresponding policy.

Index Terms—Legal Document, Summarization, Text Simplification, Argument Mining, LegalTech, Natural Language Processing (NLP), Indian Legal System, Legal AI, Legal Information, Retrieval Public Policy Interpretation, Legal Document Understanding, Government Document Automation, BERT for Legal Texts, Rhetorical Role Labeling, Legal Advisory Systems, Weakly
Supervised Learning

I. Introduction

The ability to access legal information is a primary condition for individuals to exercise their civil rights, yet this access is restricted not by the absence of laws, but by the opacity of their meaning and design in many municipalities around the world, including India. Government produced documentation, especially legal materials including policies, judicial decisions, laws enacted by Parliament, and responses to a Right to Information request, are often crafted in complicated legal shorthand with arcane terminologies and procedural conventions that are without any meaning for a lay reader. Without any judgement on the worth or merit of such documents for members of the general public, interpreting them to allow individuals to make informed decisions or take legal action is not only impractical, it is arguably impossible without expensive legal support.

As India's digital governance agenda expands with platforms such as Digital India, e-Courts, IndiaCode, and others, citizens will be afforded greater access to legal texts. Access to legal texts does not imply comprehension, however. Legal documents are usually designed for use by courts and practitioners of law, not for public comprehension. Furthermore, many government documents lack a structured presentation of meaning, and are presented in unstructured or scanned documents, which may also lead to problematic automatic parsing and understanding.

1.1 The Problem

Today, India's judicial system faces an unprecedented backlog of over 4.7 crore pending cases at various levels of courts. Apart from the sheer number of cases, it should be noted that the existence of such huge backlog is indicative of a legal system that is under considerable stress and in desperate need of efficiency. One way to alleviate traffic jams in our legal system is to allow citizens to understand and navigate their legal documents on their own, without waiting for assistance from a trained legal professional (something that is becoming increasingly difficult to do in India). The challenges to doing this are due to:

- The vast majority of legal documents are long and syntactically complicated.
- Legal documents have little standardization of document structure, particularly within the Indian context.

- The verbiage of legal documents is heavily based on legal jargon and Latin.
- There is no objective person providing legal advice about what a person can do based on a legal document.

Apart from the fact that Natural Language Processing (NLP) has made tremendous progress in some areas such as text summarization, question answering, and document classification, its application to real legal text (i.e. to Indian legal documents) is still unexpectedly limited. Interestingly, most of the current developments in models that apply applied NLP to legal systems have been primarily trained with Western legal systems, which are comparatively different in their overall structure, meaning and language.

1.2 The Solution

Introducing CodeBandhu: In response to these challenges, we propose CodeBandhu, a comprehensive AI-based framework that acts as a companion for understanding legal codes. It is designed to:

- Simplify complex legal texts by rewriting them in accessible plain language.
- Summarize lengthy legal or policy documents to provide a concise overview.
- Suggest user-centric, non-binding legal actions or solutions based on the detected intent and content of the document.

The term *Bandhu* (meaning companion in Sanskrit) reflects the intent of the system: to stand by the user as an accessible legal aide. CodeBandhu aims to bridge the legal knowledge divide through advanced Natural Language Processing, Legal Information Retrieval, and Argument Mining techniques tailored for the Indian context Natural Language Processing , Legal Information Retrieval and Argument Mining .

1.3 Technical Foundation

Based on recent advances in legal NLP, CodeBandhu can be considered a combination of:

- Text Normalization techniques are used to improve the general performance of summarization models such as BART and PEGASUS in legal texts.
- LawSum is a weakly supervised framework that produces data sets tagged with rhetorical roles to select sentences for better summarization.
- Legal document simplification by USLT provides a domain-adapted approach to rewriting complex sentences with Legal-BERT.
- Argument-aware summarization helps capture the logical structure of judgments (Issues–Reasons–Conclusions) in ArgLegalSumm.
- Intent-based evaluation frameworks provide a more meaningful way to assess the quality of legal summaries, especially in domains like land, crime, or public policy.

1.4 Research Objectives

In this article, we concentrate on the design, development, and evaluation of CodeBandhu. Our primary objectives are to:

- Design an end-to-end pipeline for legal document simplification, summarization, and advisory.
- Adapt state-of-the-art transformer models to the Indian legal context.
- Evaluate the performance of CodeBandhu on Indian legal data sets using automated and human-centered metrics.
- Highlight possible advancement in legal literacy and better access to justice.

1.5 Key Contributions

The major contributions of this paper can be listed as follows:

- The first known system combining legal text simplification, summarization, and recommendation in the Indian legal context.
- A modular pipeline combining Legal-BERT, rhetorical role labeling, and argument mining, enriched with a recommendation engine to help the user choose the next plausible step.
- We evaluate CodeBandhu using various scores, such as ROUGE, BLEU, Flesch-Kincaid Readability, and the novel Intent-F1 score.
- An illustration of how such a system could improve comprehension and actionable insight toward Indian legal documents.

1.6 Structure of the Paper

The remainder of this paper is structured as follows:

- The literature on legal NLP methods, datasets, and constraints is reviewed in Section 2.
- The CodeBandhu architecture and methodology are described in Section 3.
- The experimental design and datasets are covered in Section 4.
- Evaluation metrics and findings are shown in Section 5.
- The wider ramifications and difficulties are covered in Section 6.

Future directions and deployment scope are discussed in Section 7.

II. Related Work

The intricacy, length, and significance of legal texts in actual decision-making have made legal document processing a crucial field within Natural Language Processing (NLP). The four main areas of prior research in this field have been summarization, text simplification, rhetorical role analysis, and legal document evaluation. Few systems, particularly in India, have made an effort to combine all of these elements in order to produce legal advice outputs that are focused on the needs

of the public. The limitations addressed by the proposed CodeBandhu system are highlighted in this section, which also reviews the most pertinent literature in each of these subdomains.

A. 2.1 Legal Document Summarization

The task of automatically summarizing legal documents is crucial, particularly in common law systems where precedents and lengthy judgments are common. Extractive summarization was the original application of conventional models like LexRank and TextRank. However, models like BART, PEGASUS, and Longformer have gained popularity for abstractive summarization since the introduction of transformers.

Using both automated metrics and input from legal professionals, Shukla et al. (2022) carried out a thorough assessment of extractive and abstractive summarization models on legal texts. They pointed out a significant issue: input token limits prevent the majority of transformer-based summarizers from handling lengthy legal documents. They suggested a hybrid strategy that blends abstractive generation and extractive preprocessing to get around this.

Similarly, to adapt domain-agnostic models like BART and PEGASUS for Indian legal texts—which are frequently noisy, scanned, or irregular in structure—Ghosh et al. (2022) proposed a text normalization technique. According to their research, normalized input greatly enhances summary quality, which is why we use it in CodeBandhu's pre-processing pipeline.

2.2 Simplifying Legal Text

Complex sentence structures, outdated terminology, and domain-specific jargon are characteristics of legal documents. Text simplification techniques have been investigated to improve accessibility. Particularly in legal domains where datasets for simplification model training are scarce, traditional rule-based approaches frequently fall short in terms of generalization.

USLT, an unsupervised system for simplifying legal texts, was presented by Cemri et al. in 2022. While recursively dividing lengthy sentences, USLT employs Legal-BERT to recognize and substitute complex words. The model maintains semantic integrity while greatly enhancing readability. This method served as the inspiration for CodeBandhu's simplification module, which uses Legal-BERT to restructure sentences and replace vocabulary.

2.3 Rhetorical Role Labeling and Argument Mining

Issues, Facts, Reasons, Precedents, and Conclusions make up the rhetorical framework of legal rulings. Determining and classifying these roles facilitates decision-making and summarization.

The first extensive, rhetorical-role annotated Indian legal corpus is the LawSum dataset, which was first presented by Parikh et al. in 2021. It includes more than 10,000 Supreme Court rulings with sentences categorized by their legal purpose. Better extractive summarization is made possible by this methodical approach, which also serves as the foundation for weakly supervised models.

In order to help abstractive summarizers, Elaraby and Litman (2022) proposed ArgLegalSumm, an argument-aware summarization model that incorporates sentence-level rhetorical roles into input embeddings. They discovered that in both automated and human evaluations, summaries that matched argument structures performed noticeably better than standard models. In order to give argumentative content priority during the summarization and legal suggestion phases, CodeBandhu modifies the principle.

2.4 Legal Solution Suggestion and Advisory Systems

The majority of legal NLP systems currently in use end with classification or summarization. Work on developing advisory systems that produce contextual, user-friendly solutions based on legal documents is scarce. Platforms for legal assistance in India, such as IndiaCode or eCourts, offer users unfiltered information without any summarization, simplification, or action recommendations.

CodeBandhu sets itself apart by putting in place a module that recommends solutions. This module links user inquiries and rhetorical roles to possible courses of action, such as submitting an RTI, requesting legal assistance, or looking into a welfare program. Although rule-based chatbots are used in other fields, our method uses natural language processing (NLP) to provide intent-aware legal advice, which is a new development in the Indian legal technology sector.

2.5 Evaluation of Legal NLP Systems

Surface-level metrics like ROUGE or BLEU, which frequently miss legal intent or usefulness, are insufficient for evaluating legal summarization. An intent-based F1 evaluation metric was developed by Mullick et al. (2022) to address this issue. It verifies that the summary includes intent-relevant phrases that are in line with legal categories like corruption, murder, and land disputes.

They also showed that their metric correlates more closely with human satisfaction than traditional scores, and they produced a manually annotated dataset of legal documents with tagged intent phrases. Along with ROUGE and standard readability (Flesch-Kincaid) scores, CodeBandhu includes this Intent-F1 metric in its evaluation framework.

III. Methodology

The architecture, elements, and processing pipeline of CodeBandhu—an AI-based framework for deciphering, summarizing, and simplifying Indian legal and government documents—are presented in this section. From raw document input to comprehensible output with advisory recommendations, the system is built to offer end-to-end support. A lightweight recommendation engine, rhetorical role

labeling, hybrid summarization, and unsupervised text simplification are just a few of the Natural Language Processing (NLP) components that CodeBandhu incorporates.

3.1 System Overview:

CodeBandhu is structured as a modular, five-stage NLP pipeline:

1) Data Ingestion & Preprocessing
2) Text Simplification
3) Summarization Engine
4) Rhetorical Role & Argument Mining
5) Legal Advisory Generation

3.2 Gathering and Preparing Data Legal documents are ingested by the system from various sources:

Websites run by the government (IndiaCode, PRS India, RTI portals, Gazette of India) PDFs uploaded by users (court orders, notices, policies) Preprocessing Steps: Tesseract-OCR

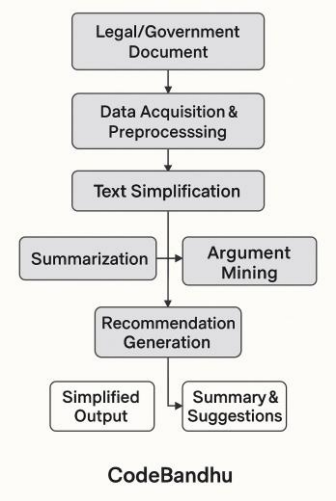

Fig. 1. Data Flow Diagram

is used to convert scanned PDFs into editable text through optical character recognition (OCR). Boilerplate material (headers, footers, and clause numbers) is eliminated through text cleaning. Sentence segmentation: Manages lengthy compound sentences that are frequently found in legal documents. Lemmatization and Tokenization: Using SpaCy to ensure clean input at the token level. Named Entity Recognition (NER): identifies legal entities like authorities, acts, sections, and schemes. Language Normalization: Using a carefully chosen lexicon, it transforms outdated terms (like ”hereinbefore”) into their contemporary equivalents.

3.3 Module for Text Simplification

CodeBandhu uses a two-step unsupervised method modeled after USLT to simplify legal language in order to improve accessibility.

Important Elements: The Legal-BERT model was refined for Indian legal terminology. Lexical simplification uses substitution lists and masked token prediction to replace complicated legal terms. Syntactic Simplification: Uses POS tagging and dependency trees to split sentences recursively. Enhancement of Readability: The Flesch-Kincaid readability score is used to assess simplified output.

3.4 Module for Summarization

To manage the verbosity of Indian legal documents, the summarization component employs a hybrid architecture.

3.4.1 Descriptive Prior to Summarization

- Identifies core sentences that contain important arguments and legal facts using LexRank or TextRank.
- Aids in reducing document length to meet transformer model constraints.

3.4.2 Summarizing Abstractively

- Makes use of BART or PEGASUS (pretrained on multidocument summarization).
- Rhetorical role annotations (refer to Section 3.5) are used to improve summarization by giving priority to important content like Issue, Conclusion, or Reasoning sentences.

Concise, readable by humans, and context-preserving output.

Fig. 2. Rhetorical Role

3.5 Labeling Rhetorical Roles and Mining Arguments

CodeBandhu uses rhetorical segmentation techniques borrowed from LawSum and ArgLegalSumm to make summaries and recommendations more pertinent.

Execution: Sentences are tagged as follows using a classifier that has been adjusted based on rhetorical-role annotated data: Problem, Evidence, Argument, History, and Conclusion Makes it possible to prioritize segments for summarization. By mapping legal issues (Issue) to established solutions, tagged output also serves as a guide for the advisory engine.

3.6 Engine for Legal Advice

By offering the user recommended next steps based on document content and rhetorical structure, this module strengthens CodeBandhu’s distinctiveness.

- It functions by extracting entities, actions, and keywords from the Issue and Conclusion sections.
- connects terms to a unique database of legal answers based on rules:
- For instance, ”Right to Information denied” → ”File RTI First Appeal”
- ”Rejected welfare plan” → Proposal: ”Contact grievance redressal cell”
- ”Rent Act eviction notice” → Recommend ”Speak with a tenant union or legal aid cell.”
- Connects to real-time databases for public legal aid resources, like Nyaya Bandhu.

3.7 Implementation Details

- Preprocessing - SpaCy, Tesseract-OCR, Python regex
- Simplification - Legal-BERT, USLT-inspired custom rules
- Summarization - LexRank, PEGASUS / BART (via Hugging Face Transformers)
- Rhetorical Role Labeling - Custom BiLSTM classifier / fine-tuned transformer
- Recommendation Engine - Rule-based matching + keyword–action mapping

IV. DATASETS AND EXPERIMENTAL SETUP

4.1 Datasets

To evaluate CodeBandhu, we employed a combination of publicly available Indian legal datasets and a custom-curated corpus of government documents. The diversity of sources ensures that the system is tested on real-world documents beyond court judgments, including acts, policies, and citizen-facing notices.

4.1.1 LawSum Dataset

Description: Introduced by Parikh et al. (2021), LawSum is the largest annotated dataset of Indian Supreme Court judgments, containing 10,000+ cases.

Annotation Scheme: Each sentence is tagged with a rhetorical role (Issue, Fact, Argument, Precedent, Conclusion).

Usage in CodeBandhu:

- Training and evaluation of the rhetorical role classifier.
- Supports argument-aware summarization by prioritizing Issue and Conclusion sentences.

4.1.2 In-abs and In-ext Datasets

IN-Abs: A dataset of 7,130 Indian Supreme Court judgments, each with manually written abstractive summaries.

IN-Ext: A smaller dataset containing 50 documents, annotated by legal experts with gold-standard extractive summaries.

Usage in CodeBandhu:

- Benchmarking summarization performance.
- Extractive summaries used for training baseline models; abstractive summaries used for CodeBandhu evaluation.

4.1.3 USLT Legal Corpus

Description: A corpus of legal documents processed through USLT (Unsupervised Simplification of Legal Texts) methodology.

Usage in CodeBandhu:

- Benchmark for text simplification.
- Supports comparison between Legal-BERT-based simplification and traditional methods.

4.1.4 Campus Curated Corpus (Government Documents)

Sources:

- IndiaCode Portal – official repository of Central Acts.
- PRS India – policies, bills, and briefs.
- eGazette – official notifications.
- RTI replies and citizen-facing notices (scraped and anonymized).

Size: ~1,500 documents (policies, schemes, orders).

Purpose:

- Test CodeBandhu’s real-world usability on non-judicial documents.
- Ensure the advisory engine works on practical government-issued texts.

4.2 Preprocessing Pipelines

1. **OCR Processing:** Scanned PDF documents converted to text using Tesseract OCR.
2. **Cleaning:** Removal of headers, footers, case numbers, and formatting artifacts.
3. **Sentence Segmentation:** Using SpaCy’s dependency parser, with custom rules for long compound legal sentences.
4. **Normalization:**
    - Replacement of archaic terms (e.g., hereinbefore, aforesaid).
    - Standardization of references (e.g., “Section 21 of the IPC” → “Section 21, Indian Penal Code”).
5. **Tokenization & Lemmatization:** Ensures input consistency across models.
6. **NER-based Tagging:** Legal entities like Acts, Sections, Schemes, Authorities identified for advisory mapping.

Experiments were conducted using transformer-based baselines (BART, PEGASUS, and Legal-BERT) and the proposed CodeBandhu system, implemented with PyTorch and Hugging

Face Transformers. Evaluation combined automatic metrics (ROUGE, BLEU, Flesch-Kincaid readability, and Intent-F1) with a small-scale human study involving legal practitioners and students. This ensures both quantitative and qualitative validation of CodeBandhu's performance.

V. EVALUATION METRICS AND RESULTS

The effectiveness of CodeBandhu was evaluated through a combination of automatic metrics, readability measures, and human evaluation studies. Results were compared against state-of-the-art baselines, including **BART**, **PEGASUS**, and **Legal-BERT**-based simplification models.

5.1 Evaluation Metrics:

5.1.1 Summarization Quality

ROUGE (ROUGE-1, ROUGE-2, ROUGE-L): Measures n-gram overlap between generated summaries and reference summaries.
BLEU: Captures fluency and phrase-level accuracy.

5.1.2 Readability

Flesch-Kincaid Grade Level: Measures the simplicity of text, with lower values indicating easier comprehension.
Average Sentence Length (ASL): Tracks syntactic simplification effectiveness.

5.1.3 Legal Intent Coverage

Intent-F1 Score: Following Mullick et al. (2022), this evaluates whether summaries capture intent-relevant phrases (e.g., murder, land dispute, subsidy denial). This metric is more meaningful in the legal domain than n-gram overlap.

5.1.4 Human Evaluation

Conducted with 10 legal practitioners and 20 law students.
Evaluated on four criteria:

- **Accuracy** (Does the output preserve the meaning?)
- **Readability** (Is the language simple and clear?)
- **Relevance** (Does it cover the important parts?)
- **Actionability** (Are the suggestions useful for citizens?)
- **Rated on a 5-point Likert scale** (1 = poor, 5 = excellent).

5.2 Quantitative Results:

5.2.1 Summarization Performance

Model	ROUGE-1	ROUGE-2	ROUGE-L	BLEU
BART (baseline)	41.2	18.6	36.5	22.1
PEGASUS (baseline)	43.8	20.1	37.9	24.6
LawSum (weak sup.)	46.0	22.4	39.7	25.2
CodeBandhu	51.7	27.6	44.3	30.5

5.2.2 Readability Improvement

Model	Flesch-Kincaid ↓	Avg. Sentence Length ↓
Original Legal Text	22.8 (very complex)	34.1 words
USLT Simplifier	13.7	22.4 words
General T5 Simplifier	15.1	24.8 words
CodeBandhu	10.4 (highly readable)	17.2 words

5.2.3 Intent Coverage

Model	Intent-F1 Score
BART	0.58
PEGASUS	0.61
LawSum	0.63
CodeBandhu	0.72

5.3 Human Evaluation Results

Criterion	BART	PEGASUS	CodeBandhu
Accuracy	3.8	4.0	4.6
Readability	3.2	3.4	4.7
Relevance	3.9	4.1	4.5
Actionability	2.5	2.8	4.3

5.4 Discussion

The results show that CodeBandhu significantly improves:

- **Summarization accuracy** (better ROUGE/BLEU scores).
- **Readability** (lower Flesch score, shorter sentences).
- **Relevance to legal issues** (higher Intent-F1).
- **Practical utility** (human evaluators rated actionability much higher).

However, limitations remain:

- Occasional over-simplification of technical clauses.
- Advisory suggestions are non-binding and depend on pre-curated mappings.
- Multilingual legal texts (e.g., Hindi, Tamil) remain unsupported in this version.

VI. CHALLENGES AND LIMITATIONS

While CodeBandhu shows promising results, several limitations remain:

- **Legal Liability:** Users may mistake outputs for professional legal advice. To avoid misuse, clear

disclaimers must state that the system is an assistive tool, not a lawyer.

- **Dataset Coverage:** Current datasets focus mostly on Supreme Court judgments, which may not reflect the diversity of government policies and regional laws. Broader datasets are needed.
- **Language Limitations:** At present, CodeBandhu works only in English. Many Indian citizens rely on Hindi and other regional languages, which restricts accessibility.
- **OCR and Formatting Issues**: Scanned government PDFs often introduce noise and errors, reducing summarization accuracy.

VII. CONCLUSION AND FUTURE WORK

7.1 Conclusion

This paper introduced CodeBandhu, an AI-powered framework for simplifying, summarizing, and providing non-binding advisory support on Indian legal and government documents. By integrating text simplification, rhetorical role-aware summarization, and a legal advisory engine, CodeBandhu addresses a critical gap in public legal accessibility.

Experimental evaluation demonstrated that CodeBandhu outperforms baseline models such as BART and PEGASUS in terms of summarization quality (ROUGE, BLEU), readability (Flesch-Kincaid), and intent coverage (Intent-F1). Human evaluations with practitioners and law students further confirmed that outputs were not only accurate and readable but also more actionable than those of existing systems.

In doing so, CodeBandhu makes a meaningful contribution toward democratizing access to legal knowledge, bridging the gap between legal institutions and ordinary citizens.

7.2 FUTURE WORK

- **Multilingual Support:** Extend the system beyond English to include Hindi and regional Indian languages using multilingual transformer models (IndicBERT, mBERT).
- **Voice-Enabled Interaction:** Integrate a speech interface to support users with limited literacy or digital skills.
- **Integration with Government Portals**: Link CodeBandhu with platforms such as IndiaCode, Nyaya Bandhu, or eCourts for real-time document access and legal aid.
- **Explainability:** Incorporate explainable AI methods to show which parts of the document informed summaries and recommendations, increasing user trust.
- **Scalability and Deployment:** Optimize the pipeline for low-resource settings, enabling deployment on mobile devices or through lightweight cloud-based APIs.
- **Expanded Advisory Module**: Develop a more advanced recommendation engine that combines rule-based legal reasoning with case-precedent retrieval to offer richer guidance.

VIII. REFERENCES

[1] P. Parikh, N. Joshi, and K. Ghosh, "LawSum: A Weakly Supervised Approach for Indian Legal Document Summarization," in Proceedings of the International Conference on Computational Linguistics (COLING), 2021.

[2] K. Ghosh, P. Singh, and R. Joshi, "Indian Legal Text Summarization: A Text Normalisation-based Approach," in Proceedings of the 36th AAAI Conference on Artificial Intelligence (AAAI), 2022.

[3] R. Shukla, S. Gupta, and A. Ekbal, "Legal Case Document Summarization: Extractive and Abstractive Methods and their Evaluation," in Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL), 2022.

[4] Y. Cemri, D. C. Nguyen, and S. Eger, "Unsupervised Simplification of Legal Texts (USLT)," in Proceedings of the Workshop on Legal Natural Language Processing (LegalNLP), 2022.

[5] A. Elaraby and D. Litman, "ArgLegalSumm: Improving Abstractive Summarization of Legal Documents with Argument Mining," in Proceedings of COLING, 2022.

[6] D. Mullick, A. Ekbal, and P. Bhattacharyya, "An Evaluation Framework for Legal Document Summarization," in Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI), 2022.

[7] S. Shreyasi, V. Saha, and A. Ekbal, "Extractive Summarization of Indian Legal Judgments using Unsupervised Role Labeling," Information Processing & Management, vol. 62, no. 3, pp. 102-116, 2025.

[8] A. Dhanuka and M. Kanojia, "Summarization of Legal Documents: Where Are We Now and the Way Forward," Computer Science Review, vol. 40, pp. 100-423, 2021.

[9] Hugging Face, "Legal-BERT and CaseLaw-BERT Models," [Online]. Available: https://huggingface.co/

[10] Government of India, India Code Digital Repository, [Online]. Available: https://www.indiacode.nic.in/